

HOWTOCFD23

Introduction to OpenMP

Instructor: Niranjana S. Ghaisas

OpenMP

- Standard for shared-memory parallel programming
- Anybody can implement these standards: gcc, intel, PGI, ...
- Detailed Specifications:
 - <https://www.openmp.org/>
- Extensions to high-level languages (e.g. C, Fortran, C++)
- High-level 'wrapper' to an internal implementation of **threads**, e.g. POSIX threads
- Practically: add a set of **compiler directives** to a serial code

OpenMP

- Standard for shared-memory parallel programming
- Anybody can implement these standards: gcc, intel, PGI, ...
- Extensions to high-level languages (e.g. C, Fortran, C++)
- History:

OpenMP Version	Year of Release	GNU C Support	Intel C Support
3.0	2008	gcc 4.4.0 +	icc 12.0+
4.0	2013	gcc 4.9.1 +	icc 15.0 +
4.5	2015	gcc 7.1 +	icc 17.0 +
5.0	2019	gcc 9.1 +	icc 19.1 +

OpenMP

- Threads:
 - A strand of execution within a *process*
 - Process: collection of instructions + memory (stack and heap)
 - Threads: parts of the process with same heap memory but individual stack memory
 - A process can have one or multiple threads
- Compiler Directives:
 - Special instructions to compiler (sometimes the preprocessor, e.g. `#define`, `#include`)
 - Specific directive used in OpenMP is
`#pragma omp parallel`

OpenMP: Workflow

```
int main(){
    // C code starts here

    int i;
    .....
    {
        // one thread working
        everywhere
        .....
    }
    .....
    return 0;
}
```

Program
Execution



```
int main(){
    // C code starts here

    int i;
    ....
    #pragma omp parallel {
        // multiple threads
        working here
        ....
    }
    ....
    return 0;
}
```

OpenMP: Workflow

```
int main(){
    // C code starts here

    int i;
    .....
    {
        // one thread working
        everywhere
        .....
    }
    .....
    return 0;
}
```

- gcc myprog.c
- ./a.out

Program
Execution



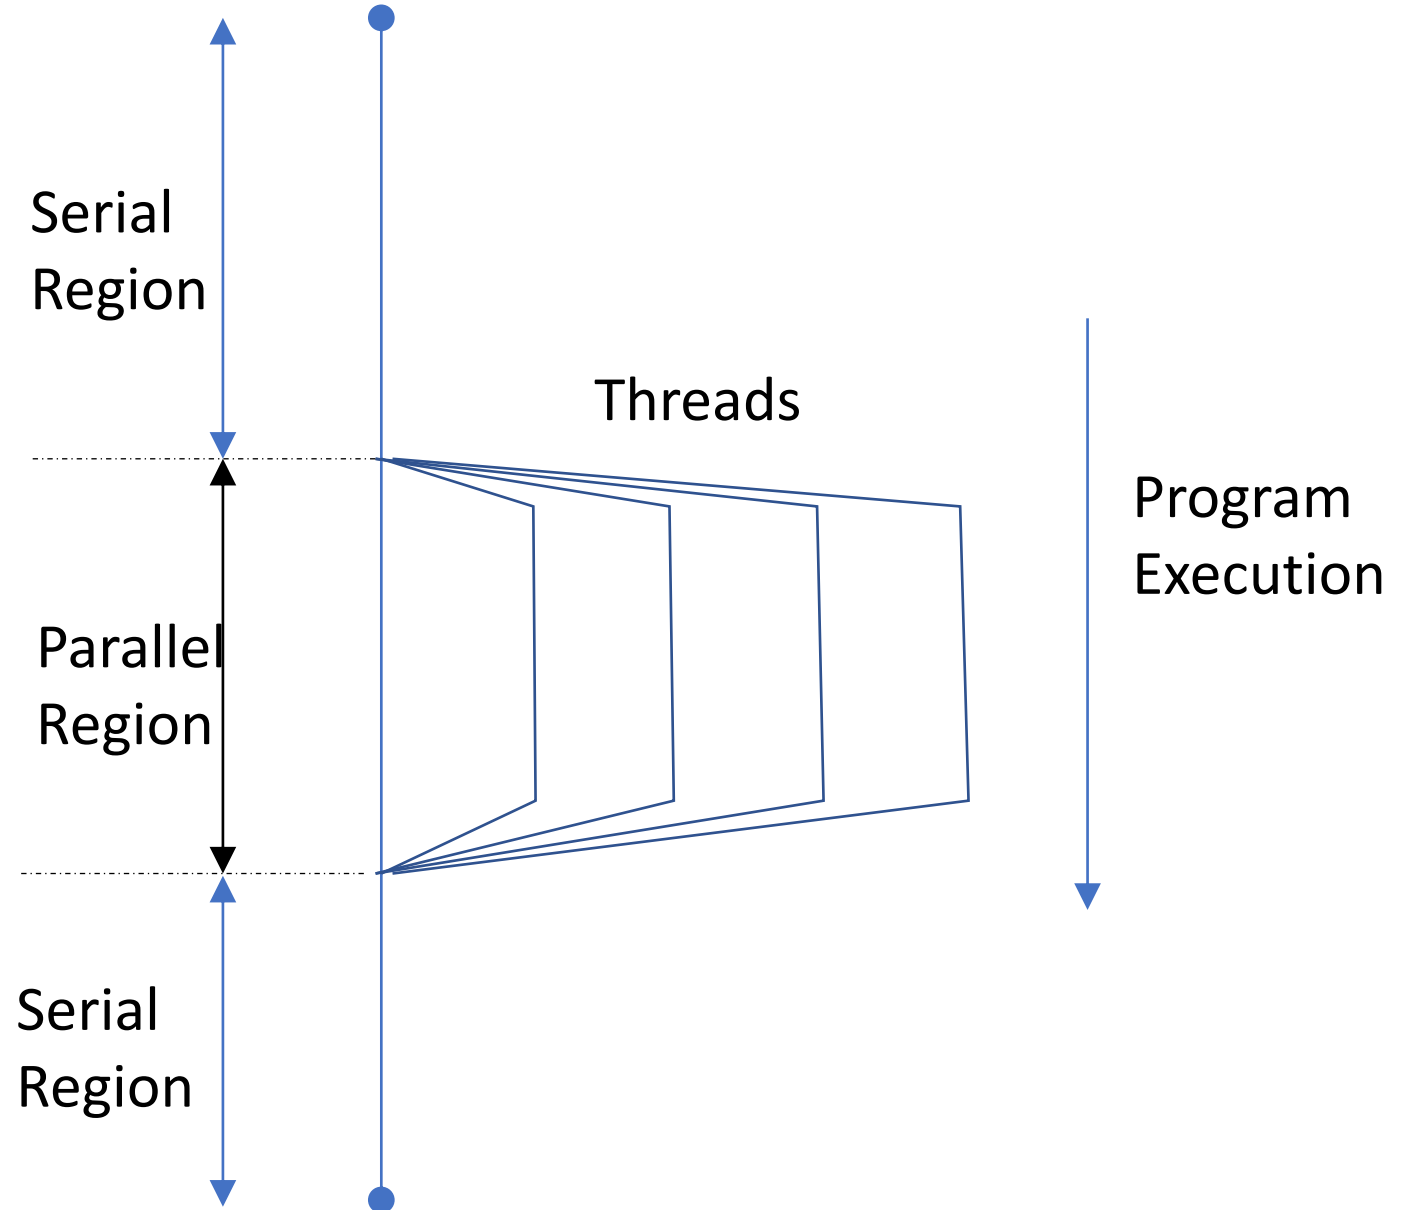
```
int main(){
    // C code starts here

    int i;
    ....
    #pragma omp parallel {
        // multiple threads
        working here
        ....
    }
    ....
    return 0;
}
```

- gcc -fopenmp myprog.c
- ./a.out


OpenMP: Workflow

```
int main(){  
    // C code starts here  
  
    int i;  
    ....  
    #pragma omp parallel {  
        // multiple threads  
        working here  
        ....  
    }  
    ....  
    return 0;  
}
```



OpenMP: Fork/Join Model

```
int main(){  
    // C code starts here  
  
    int i;  
    ....  
    #pragma omp parallel {  
        // multiple threads  
        working here  
        ....  
    }  
    ....  
    return 0;  
}
```



OpenMP: Construct and Clauses

```
int main(){
    // C code starts here

    int i;
    ....
    #pragma omp parallel {
        // multiple threads
        working here
        ....
    }
    ....
    return 0;
}
```

- #pragma omp **construct** [**clause(options)**]
- **construct**: main instruction that applies to the following block of code
- **clause(options)**: optional, some default values assumed if not specified explicitly

OpenMP: parallel construct

```
int main(){
    // C code starts here

    int i;
    ....
    #pragma omp parallel {
        // multiple threads
        working here
        ....
    }
    ....
    return 0;
}
```

- #pragma omp **construct** [**clause(options)**]
- **construct**: main instruction that applies to the following block of code
- **clause(options)**: optional, some default values assumed if not specified explicitly

Number of Threads

- Three ways of setting number of threads:
- Environment variable `OMP_NUM_THREADS`
- Library Routine: `omp_set_num_threads()`
- `num_threads` clause

Linux:

- `gcc -fopenmp myprog.c`
- `export OMP_NUM_THREADS = 8`
- `./a.out`

Number of Threads

- Three ways of setting number of threads:
- Environment variable OMP_NUM_THREADS
- Library Routine: `omp_set_num_threads()`
- `num_threads` clause

```
int main(){
    // C code starts here
    int p = 8;
    omp_set_num_threads(p);
    ... // master thread only
    #pragma omp parallel {
        // p threads working here
        ...
    }
    ... // master thread only
    return 0;
}
```

Number of Threads

- Three ways of setting number of threads:
- Environment variable
OMP_NUM_THREADS
- Library Routine:
omp_set_num_threads()
- `num_threads` clause

```
int main(){
    // C code starts here

    int p = 8;
    ...
    #pragma omp parallel num_threads (p)
    {
        // multiple threads working here
        ...
    }
    ...
    return 0;
}
```

Number of Threads

- Three ways of setting number of threads:
- Environment variable OMP_NUM_THREADS
- Library Routine: `omp_set_num_threads()`
- `num_threads` clause
- Hello world example

Number of Threads

- Three ways of setting number of threads:
- Environment variable OMP_NUM_THREADS
- Library Routine: `omp_set_num_threads()`
- `num_threads` clause
- Hello world example


- Question: If more than one of the above is present?
 - Experiment for yourself and find out.
- Question: How many threads can you run?
 - Depends on the “system”. Effectively, no limit.
- Question: How many *should* I run?
 - Always experiment. Thumb rule: number of cores

Distinguishing Between Threads

```
int main(){
    int n = 8;
    omp_set_num_threads(n);
    ... // master thread only

    #pragma omp parallel {
        int tid = omp_get_thread_num();
        printf("Hello world %d", tid);
    }
    ... // master thread only
    return 0;
}
```

- Library Routine: `omp_get_num_threads()`



OpenMP Basics Summary/Overview

Purpose	construct/clause
Create threads	<ul style="list-style-type: none">• <code>parallel</code> construct
Distribute work	<ul style="list-style-type: none">• <code>for</code> construct• <code>sections</code> construct
Variable scoping	<ul style="list-style-type: none">• <code>shared</code>, <code>private</code>, <code>firstprivate</code>, <code>lastprivate</code>, <code>reduction</code>, <code>collapse</code>, ... clauses
Synchronization	<ul style="list-style-type: none">• <code>critical</code>, <code>atomic</code>, <code>barrier</code>, <code>nowait</code> constructs
Functions/environment variables	<ul style="list-style-type: none">• <code>omp_set_num_threads</code>, <code>omp_get_thread_num</code>• <code>OMP_NUM_THREADS</code>, <code>OMP_SCHEDULE</code>

OpenMP Parallelism: Other Languages

- Exactly analogous support for Fortran (replace `#pragma` with `!$OMP`)
- Some functionality supported in Matlab as well (e.g. `parfor`)
- Python: the picture is complicated
 - Python operates under 'GIL' (Global Interpreter Lock); only one thread can be active at a time
 - any compiled code (C/Fortran/...) called from within Python can be OpenMP-parallel

OpenMP Basics Summary/Overview

Purpose	construct/clause
Create threads	<ul style="list-style-type: none">• <code>parallel</code> construct
Distribute work	<ul style="list-style-type: none">• <code>for</code> construct• <code>sections</code> construct
Variable scoping	<ul style="list-style-type: none">• <code>shared</code>, <code>private</code>, <code>firstprivate</code>, <code>lastprivate</code>, <code>reduction</code>, <code>collapse</code>, ... clauses
Synchronization	<ul style="list-style-type: none">• <code>critical</code>, <code>atomic</code>, <code>barrier</code>, <code>nowait</code> constructs
Functions/environment variables	<ul style="list-style-type: none">• <code>omp_set_num_threads</code>, <code>omp_get_thread_num</code>• <code>OMP_NUM_THREADS</code>, <code>OMP_SCHEDULE</code>

OpenMP Basics Summary/Overview

Purpose

Create threads

Distribute work

Variable scoping

Synchronization

Functions/environment
variables

construct/clause

- `parallel` construct

- `for` construct
- `sections` construct

- `shared`, `private`, `firstprivate`,
`lastprivate`, `reduction`, `collapse`, ... clauses

- `critical`, `atomic`, `barrier`, `nowait`
constructs

- `omp_set_num_threads`, `omp_get_thread_num`
- `OMP_NUM_THREADS`, `OMP_SCHEDULE`

Sharing Work Between Threads

- Two ways:

- `for`
- `sections`

```
...  
#pragma omp parallel {  
    #pragma omp for  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
    }  
}  
...  
}
```

Split loop iterations among
available threads


Wait until all threads arrive
here (implicit barrier)

Sharing Work Between Threads

- Two ways:

- **for**
- sections

```
...  
#pragma omp parallel {  
    #pragma omp for  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
    }  
}  
...
```



```
...  
#pragma omp parallel for  
{  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
    }  
}  
...
```

Sharing Work Between Threads

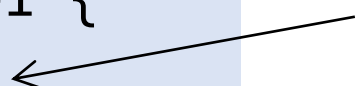
- Two ways:

- **for**

- sections

```
...  
#pragma omp parallel {  
  //#pragma omp for  
  for(i=0; i<N; i++)  
  {  
    a[i] = b[i] + c[i];  
  }  
}  
...
```

What happens if you miss the
for?



Sharing Work Between Threads

- Two ways:

- for

- sections

```
...  
#pragma omp parallel {  
    #pragma omp sections {  
        #pragma section {  
            // task 1      }  
        #pragma section {  
            // task 2      }  
        .....  
        #pragma section {  
            // task n      }  
    }  
}  
...
```

Assign one section to one thread

Wait until all threads arrive here (implicit barrier)

Sharing Work Between Threads

- Two ways:

- for

- sections

```
...
#pragma omp parallel {
    #pragma omp sections {
        #pragma section {
            // task 1      }
        #pragma section {
            // task 2      }
        .....
        #pragma section {
            // task n      }
    }
}
...
```

- Question: Which thread gets which section?

- Question: What if num_threads > num_sections?

- Question: What if num_threads < num_sections?

Sharing Work Between Threads

- Two ways:

- for

- sections

```
#pragma omp parallel {  
    ...  
    #pragma omp sections {  
        #pragma section {  
            // task 1      }  
        #pragma section {  
            // task 2      }  
        .....  
        #pragma section {  
            // task n      }  
    }  
    ...  
}
```

- Question: Which thread gets which section?
 - First come first served
- Question: What if num_threads > num_sections?
 - Latecomers remain idle
- Question: What if num_threads < num_sections?
 - Round robin assignment of sections as threads become available

schedule Clause

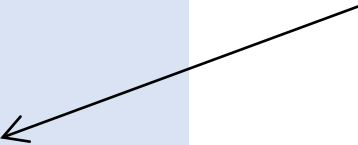
- Two ways:

- **for**

- sections

```
...  
#pragma omp parallel {  
    #pragma omp for  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
    }  
}  
...
```

Split loop iterations among
available threads



- Question: How can we control the 'split'?
- Answer: schedule clause

schedule Clause

```
...  
#pragma omp parallel {  
    #pragma omp for schedule(kind, chunksize)  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
        d[i] = do_work(i, a[i]);  
    }  
}  
...
```

- kind can be:
 - static
 - dynamic
 - guided
 - runtime
- chunksize is optional

schedule Clause

```
...  
#pragma omp parallel {  
    #pragma omp for schedule(kind, chunksize)  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
        d[i] = do_work(i, a[i]);  
    }  
}  
...
```

- kind can be:
 - **static**
 - dynamic
 - guided
 - runtime
- chunksize is optional
- Round-robin assignment of chunks
- Default chunksize = $N/\text{num_threads}$ (roughly)

schedule Clause

```
...  
#pragma omp parallel {  
    #pragma omp for schedule(kind, chunksize)  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
        d[i] = do_work(i, a[i]);  
    }  
}  
...
```

- kind can be:
 - static
 - **dynamic**
 - guided
 - runtime
- chunksize is optional
- Chunks assigned as threads become available
- Default chunksize = 1

schedule Clause

```
...  
#pragma omp parallel {  
    #pragma omp for schedule(kind, chunksize)  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
        d[i] = do_work(i, a[i]);  
    }  
}  
...
```

- kind can be:
 - static
 - dynamic
 - guided
 - runtime
- chunksize is optional
- dynamic with dynamic size of chunks (proportional to remaining work).
- chunksize sets the minimum size except for the last assignment

schedule Clause

```
...
#pragma omp parallel {
    #pragma omp for schedule(kind, chunksize)
    for(i=0; i<N; i++)
    {
        a[i] = b[i] + c[i];
        d[i] = do_work(i, a[i]);
    }
}
...
```

- kind can be:
 - static
 - dynamic
 - guided
 - runtime
- chunksize is optional
- kind and chunksize set at runtime. Details vary by implementation (gnu vs intel vs ...)

schedule Clause

```
...  
#pragma omp parallel {  
    #pragma omp for schedule(kind, chunksize)  
    for(i=0; i<N; i++)  
    {  
        a[i] = b[i] + c[i];  
        d[i] = do_work(i, a[i]);  
    }  
}  
...
```

schedule example

- kind can be:
 - static
 - dynamic
 - guided
- chunksize is optional
- When to use what?
 - Simple, predictable, roughly equal work for each iteration -> static
 - Unpredictable, highly variable work in different iterations -> dynamic, guided

OpenMP Basics Summary/Overview

Purpose

Create threads

Distribute work

Variable scoping

Synchronization

Functions/environment
variables

construct/clause

- `parallel` construct

- `for` construct
- `sections` construct

- `shared`, `private`, `firstprivate`,
`lastprivate`, `reduction`, `collapse`, ... clauses

- `critical`, `atomic`, `barrier`, `nowait`
constructs

- `omp_set_num_threads`, `omp_get_thread_num`
- `OMP_NUM_THREADS`, `OMP_SCHEDULE`

Variable Sharing

```
int main(){
    int p = 8, i;
    omp_set_num_threads(p);
    ... // master thread only


    #pragma omp parallel {
        int tid = omp_get_thread_num();
        printf("Hello world %d", tid);
    }
    ... // master thread only
    return 0;
}
```

vs

```
int main(){
    int p = 8, i, tid;
    omp_set_num_threads(p);
    ... // master thread only


    #pragma omp parallel {
        tid = omp_get_thread_num();
        printf("Hello world %d", tid);
    }
    ... // master thread only
    return 0;
}
```

Variable Sharing

```
int main(){  
    int p = 8, i;   
    omp_set_num_threads(p);  
    ... // master thread only  
  
    #pragma omp parallel {  
        int tid = omp_get_thread_num();  
        printf("Hello world %d", tid);  
    }  
    ... // master thread only  
    return 0;  
}
```

tid is private to each thread (stack)

vs

```
int main(){  
    int p = 8, i, tid;   
    omp_set_num_threads(p);  
    ... // master thread only  
  
    #pragma omp parallel {  
        tid = omp_get_thread_num();  
        printf("Hello world %d", tid);  
    }  
    ... // master thread only  
    return 0;  
}
```

tid is shared among all threads (heap)

Variable Sharing

```
int main(){
    int p = 8, i, N = 64;
    omp_set_num_threads(p);

    #pragma omp parallel for {
        for(i=0; i<N; i++) {
            a[i] = b[i] + c[i];
        }
    }
    return 0;
}
```

vs

```
int main(){
    int p = 8, i, N = 64, k;
    omp_set_num_threads(p);

    #pragma omp parallel for {
        for(i=0; i<N; i++) {
            k = i;
            a[k] = b[k] + c[k];
        }
    }
    return 0;
}
```

Variable Sharing

```
int main(){  
    int p = 8, i, N = 64;  
    omp_set_num_threads(p);  
  
    #pragma omp parallel for {  
        for(i=0; i<N; i++) {  
            a[i] = b[i] + c[i];  
        }  
    }  
    return 0;  
}
```



i is private to each thread (stack)

vs

```
int main(){  
    int p = 8, i, N = 64, k;  
    omp_set_num_threads(p);  
  
    #pragma omp parallel for {  
        for(i=0; i<N; i++) {  
            k = i;  
            a[k] = b[k] + c[k];  
        }  
    }  
    return 0;  
}
```




k is shared among all threads (heap)

Variable Sharing Rules

- Variables that are shared by default:
 - Allocated on heap
 - Declared outside the scope of the parallel construct
 - `static` or constant variables
- Variables that are private by default:
 - Declared within the scope of a `parallel` construct
 - Index variable in a `parallel for` construct
- Not paying attention to variable scoping is the biggest cause for “data race” or “race conditions”
- Debugging these issues is, very often, the most time-consuming part of OpenMP programming!

Variable Sharing: Overriding the Defaults


```
int main(){  
    int p = 8, tid;  
    omp_set_num_threads(p);  
    ... // master thread only  
  
    #pragma omp parallel {  
        tid = omp_get_thread_num();  
        printf("Hello world %d", tid);  
    }  
    ... // master thread only  
    return 0;  
}
```



tid is shared among all threads (heap)

Variable Sharing: Overriding the Defaults


```
int main(){  
    int p = 8, tid;  
    omp_set_num_threads(p);  
    ... // master thread only  
  
    #pragma omp parallel private(tid){  
        tid = omp_get_thread_num();  
        printf("Hello world %d", tid);  
    }  
    ... // master thread only  
    return 0;  
}
```



tid is private to each thread (stack)

vs

```
int main(){  
    int p = 8, tid;  
    omp_set_num_threads(p);  
    ... // master thread only  
  
    #pragma omp parallel {  
        tid = omp_get_thread_num();  
        printf("Hello world %d", tid);  
    }  
    ... // master thread only  
    return 0;  
}
```



tid is shared among all threads (heap)

Variable Sharing: Overriding the Defaults

```
int main(){
    int p = 8, tid;
    omp_set_num_threads(p);
    ... // master thread only

    #pragma omp parallel private(tid){
        tid = omp_get_thread_num();
        printf("Hello world %d", tid);
    }
    ... // master thread only
    return 0;
}
```

vs

- p-1 additional instances of tid are created (copies on stack of each thread)
- May or may not be initialized (implementation dependent, OpenMP standard does not specify initialization)

tid is private to each thread (stack)

Variable Sharing: Overriding the Defaults

```
int main(){
    int p = 8, tid;
    omp_set_num_threads(p);
    ... // master thread only

    #pragma omp parallel private(tid){
        tid = omp_get_thread_num();
        printf("Hello world %d", tid);
    }
    printf("Outside parallel %d", tid);
    return 0;
}
```

vs

- p-1 additional instances of tid are created (copies on stack)
- May or may not be initialized (implementation dependent, OpenMP standard does not specify initialization)
- What happens at the end of the parallel construct?
- Only one instance (master) remains

tid is private to each threads (stack)

Variable Sharing: Overriding the Defaults

```
int main(){
    int p = 8, tid;
    omp_set_num_threads(p);
    ... // master thread only

    #pragma omp parallel private(tid){
        tid = omp_get_thread_num();
        printf("Hello world %d", tid);
    }
    printf("Outside parallel %d", tid);
    return 0;
}
```

vs

- p-1 additional instances of tid are created (copies on stack)
- May or may not be initialized (implementation dependent, OpenMP standard does not specify initialization)
- What happens at the end of the parallel construct?
- Only one instance (master) remains
- Ways to control initialization and termination of parallel construct

tid is private to each threads (stack)

parallel construct clauses, options (partial)

```
...  
#pragma omp parallel  
    private (var1, var2, var3, var4, ...)  
    firstprivate (var1, var2, ...)  
    reduction(operator:list)  
    {  
        // parallel block  
    }  
...
```

- Control for initialization:
 - firstprivate: initialize value of var for all threads to that before parallel
- Control at termination:
 - reduction: apply operator (e.g. sum, min, max,...) on each variable in the list. Value of variable after the parallel construct is determined by this operation

parallel construct clauses, options (complete)

```
...
#pragma omp parallel
    num_threads(p)
    private (var1, var2, var3, var4, ...)
    firstprivate (var1, var2, ...)
    reduction(operator:list)
    if (boolean)
    default(shared OR none)
    shared (var5, var6)
    copyin (var1, var3, ...)
    proc_bind(...)
    allocate(...)
{ // parallel block }
...
```

- if: whether to execute in parallel
- default: for variables declared within scope
- shared: force variables to be shared across threads
- copyin: initialize variables to values existing prior to parallel (similar to firstprivate here)
- proc_bind, allocate: look up yourself

sections construct clauses, options (complete)

```
#pragma omp parallel
#pragma omp sections
    private (var1, var2, var3, var4, ...)
    firstprivate (var1, var2, ...)
    lastprivate (var3, var4, ...)
    reduction(operator:list)
    nowait
    allocate(...)
{ // parallel sections block }
...
```

- private: variable specific to each thread
- firstprivate: assign value of var1, var2 from before parallel region at the start of sections
- lastprivate: assign value from the last section to var3, var4 at the end of sections
- reduction: apply operator (e.g. sum, min, max,...) on each variable in the list. Value of variable after the sections construct is determined by this operation
- nowait: coming up
- allocate: look up yourself

for construct clauses, options (complete)

```
#pragma omp parallel
#pragma omp for
    private (var1, var2, var3, var4, ...)
    firstprivate (var1, var2, ...)
    lastprivate (var3, var4, ...)
    reduction(operator:list)
    schedule (kind, chunk_size)
    linear (var5:2, list[:linear-step])
    collapse(n)
    ordered [(n)]
    order (concurrent)
    nowait
    allocate(...)
{ // parallel block }
...
```

- lastprivate: assign value from the last iteration to variable at the end of for
- linear: increment var5 by 2 at every iteration; retain value at end
- collapse: nested for loops, n loops are combined and then distributed
- nowait: coming up
- ordered, order, allocate: look up yourself

parallel/for/sections construct examples

- private, firstprivate, example
- reduction example

OpenMP Basics Summary/Overview

Purpose

Create threads

Distribute work

Variable scoping

Synchronization

Functions/environment
variables

construct/clause

- `parallel` construct

- `for` construct
- `sections` construct

- `shared`, `private`, `firstprivate`,
`lastprivate`, `reduction`, `collapse`, ... clauses

- `critical`, `atomic`, `barrier`, `nowait`,
`single`, `master` constructs

- `omp_set_num_threads`, `omp_get_thread_num`
- `OMP_NUM_THREADS`, `OMP_SCHEDULE`

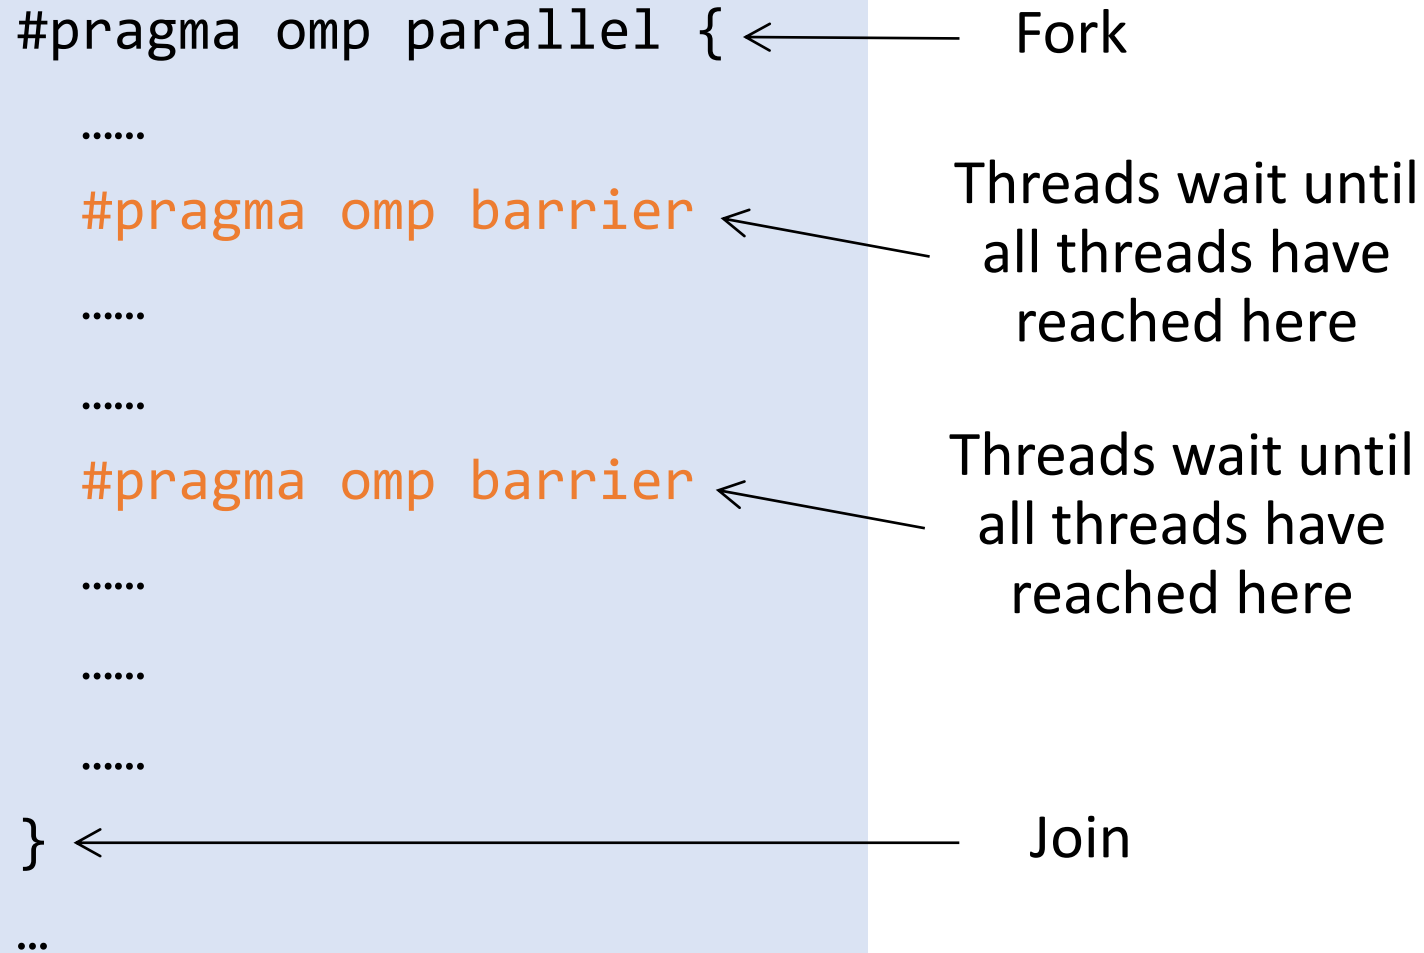
Synchronization

- No synchronization between threads unless explicitly required by the code (MIMD model)
- (Implicit) Barrier/nowait
- single/master
- critical/atomic
- ordered
- locks

```
#pragma omp parallel {  
    ..... // parallel work  
    ..... // parallel work  
    ..... // parallel work  
    ..... // synchronization  
    ..... // parallel work  
    ..... // parallel work  
    ..... // synchronization  
}
```

Synchronization: `barrier`

```
#pragma omp parallel { ← Fork
.....
#pragma omp barrier ← Threads wait until
.....                all threads have
.....                reached here
#pragma omp barrier ← Threads wait until
.....                all threads have
.....                reached here
.....
.....
.....
} ← Join
...
```



Synchronization (implicit barrier)

```
#pragma omp parallel { ← Fork
.....
#pragma omp for
for(i=0; i<N; i++) {
    a[i] = b[i] + c[i];
} ← Implicit Barrier
.....
#pragma omp for
for(i=0; i<N; i++) {
    a[i] = b[i] + c[i];
} ← Implicit Barrier
.....
} ← Join
```

(De-) synchronization: `nowait`

```
#pragma omp parallel { ← Fork
```

.....

```
#pragma omp for  
for(i=0; i<N; i++) {  
    a[i] = b[i] + c[i];  
} ←
```

Implicit
Barrier

x

.....

```
#pragma omp for  
for(i=0; i<2*N; i++) {  
    d[i] = e[i] + f[i];  
} ←
```

Implicit
Barrier

.....

```
} ← Join  
(Barrier)
```

```
#pragma omp parallel {
```

.....

```
#pragma omp for nowait  
for(i=0; i<N; i++) {  
    a[i] = b[i] + c[i];  
} →
```

.....

```
#pragma omp for  
for(i=0; i<2*N; i++) {  
    d[i] = e[i] + f[i];  
} →
```

.....

```
}
```

(De-) synchronization: `nowait`

```
#pragma omp parallel { ← Fork
```

.....

```
#pragma omp for  
for(i=0; i<N; i++) {  
    a[i] = b[i] + c[i];  
} ←
```

Implicit
Barrier

✗

```
#pragma omp parallel {
```

.....

```
#pragma omp for nowait  
for(i=0; i<N; i++) {  
    a[i] = b[i] + c[i];  
} →
```

.....

```
#pragma omp for  
for(i=0; i<2*N; i++) {  
    d[i] = e[i] + f[i];  
} ←
```

Implicit
Barrier

✗

.....

```
#pragma omp for nowait  
for(i=0; i<2*N; i++) {  
    d[i] = e[i] + f[i];  
} →
```

.....

```
} ← Join  
(Barrier)
```

.....

```
}
```

(De-) synchronization: `single/master`

```
#pragma omp parallel {  
    .....  
    #pragma omp for  
    for(i=0; i<N; i++) {  
        a[i] = b[i] + c[i];  
    }  
    #pragma omp single  
        printf("Finished loop\n");  
    .....  
    .....  
}
```

First thread to reach `single` executes; others proceed to the next line

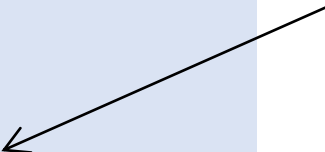
```
#pragma omp parallel {  
    .....  
    #pragma omp for  
    for(i=0; i<N; i++) {  
        a[i] = b[i] + c[i];  
    }  
    #pragma omp master  
        printf("Finished loop\n");  
    .....  
    .....  
}
```

Only master executes; others proceed to the next line

Synchronization: critical/atomic

```
#pragma omp parallel
private(lsum) shared (sum) {
    #pragma omp single
        sum = 0.0;
    lsum = 0.0;
    #pragma omp for
    for(i=0; i<N; i++)
        lsum += a[i];
    #pragma omp critical
        sum += lsum;

    .....
}
```

- Only one thread executes at a time.
 - No particular order is enforced
- 

Synchronization: critical/atomic

```
#pragma omp parallel
private(lsum) shared (sum) {
    #pragma omp single
        sum = 0.0;
    lsum = 0.0;
    #pragma omp for
    for(i=0; i<N; i++)
        lsum += a[i];
    #pragma omp critical
        sum += lsum;

    .....
}
```

- Only one thread executes at a time.
- No particular order is enforced

```
#pragma omp parallel
private(lsum) shared (sum) {
    #pragma omp single
        sum = 0.0;
    lsum = 0.0;
    #pragma omp for
    for(i=0; i<N; i++)
        lsum += a[i];
    #pragma omp atomic
        sum += lsum;

    .....
}
```

Synchronization: `critical`/`atomic` differences

```
#pragma omp parallel {  
.....  
#pragma omp critical  
    sum += lsum;  
.....  
.....  
#pragma omp critical {  
    .... // Any operation  
    .... // Any operation  
}  
.....  
}
```

Critical:
One or more
operations. No
restrictions

Atomic:
Only one
operation; Only
specific operations
(assignment)

```
#pragma omp parallel {  
.....  
#pragma omp atomic  
    sum += lsum; ✓  
.....  
.....  
#pragma omp atomic { ✗  
    ....  
    ....  
}  
.....  
}
```

Synchronization: `critical/atomic`

```
#pragma omp parallel {  
    .....  
    #pragma omp critical  
        sum += lsum;  
    .....  
    .....  
    #pragma omp critical {  
        if(current_val > max_val)  
            max_val = current_val;  
    }  
    .....  
}
```

- If more than one critical section occurs, they are mutually exclusive
 - Only one thread at a time across all critical sections
- Synchronization is being forced across critical sections; can affect performance

Synchronization: `critical/atomic`

```
#pragma omp parallel {  
    .....  
    #pragma omp critical (summ)  
        sum += lsum;  
    .....  
    .....  
    #pragma omp critical (maxx) {  
        if(current_val > max_val)  
            max_val = current_val;  
    }  
    .....  
}
```

- If more than one critical section occurs, they are mutually exclusive
 - Only one thread at a time across all `unnamed` critical sections
- Synchronization is being forced across critical sections; can affect performance
- Remedy: named critical sections
 - Named critical sections are not mutually exclusive

Synchronization: ordered

```
#pragma omp parallel {  
    .....  
    #pragma omp for ordered  
    for(i=0; i<N; i++) {  
        .....  
        .....  
        #pragma omp ordered  
            printf("Var value on %d =  
%f\n", tid, val);  
    }  
    .....  
}
```

- #pragma omp **construct** [**clause(options)**]
- ordered clause: force a block in a for loop to be executed as if it were serial
- ordered construct: indicates that a ordered clause appears somewhere in this construct

OpenMP Basics Summary/Overview

Purpose	construct/clause
Create threads	<ul style="list-style-type: none">• <code>parallel</code> construct
Distribute work	<ul style="list-style-type: none">• <code>for</code> construct• <code>sections</code> construct
Variable scoping	<ul style="list-style-type: none">• <code>shared</code>, <code>private</code>, <code>firstprivate</code>, <code>lastprivate</code>, <code>reduction</code>, <code>collapse</code>, ... clauses
Synchronization	<ul style="list-style-type: none">• <code>critical</code>, <code>atomic</code>, <code>barrier</code>, <code>nowait</code> constructs
Functions/environment variables	<ul style="list-style-type: none">• <code>omp_set_num_threads</code>, <code>omp_get_thread_num</code>• <code>OMP_NUM_THREADS</code>, <code>OMP_SCHEDULE</code>

More Examples

More Examples

```
.....  
#pragma omp parallel {  
    ..... //{code A}  
  
    #pragma omp for  
    for(i=0; i<N; i++) {  
        ..... //{code B}  
    }  
  
    ..... // {code C}  
}  
.....
```

Fork

Barrier


Join

tid = 0

1

2

3



More Examples

.....

```
#pragma omp parallel {
```

```
..... //{code A}
```

```
#pragma omp for
```

```
for(i=0; i<N; i++) {
```

```
..... //{code B}
```

```
}
```

```
..... // {code C}
```

```
}
```

.....

Fork

tid = 0

1

2

3

Code A

Code A

Code A

Code A

[0, N/4]

[N/4+1
,N/2]

[N/2+1,
3N/4]

[3N/4
+1, N]

Barrier

Code C

Code C

Code C

Code C

Join

More Examples

.....

```
#pragma omp parallel {
```

```
..... //{code A}
```

```
#pragma omp sections {
```

```
#pragma section {
```

```
..... //{code B}
```

```
}
```

```
#pragma section {
```

```
..... //{code C}
```

```
}
```

```
}
```

```
..... //{code D}
```

```
}
```

.....

Fork

Barrier


Join

tid = 0

1

2

3



More Examples

.....

```
#pragma omp parallel {
```

```
..... //{code A}
```

```
#pragma omp sections {
```

```
#pragma section {
```

```
..... //{code B}
```

```
}
```

```
#pragma section {
```

```
..... //{code C}
```

```
}
```

```
}
```

```
..... //{code D}
```

```
}
```

.....

Fork

Barrier


Join

tid = 0

1

2

3



More Examples

.....

```
#pragma omp parallel {
```

```
..... //{code A}
```

```
#pragma omp sections {
```

```
#pragma section {
```

```
..... //{code B}
```

```
}
```

```
#pragma section {
```

```
..... //{code C}
```

```
}
```

```
}
```

```
..... //{code D}
```

```
}
```

.....

Fork

Barrier

Join

tid = 0

1

2

3

Code A

Code A

Code A

Code A

Code B

Code C

Code D

Code D

Code D

Code D

Example: Find the bug

```
/* Declare and initialize a, b, c, chunk, ... */

#pragma omp parallel for shared(a,b,c,chunk)
private(i,tid) schedule(static,chunk) {
    tid = omp_get_thread_num();
    for (i=0; i < N; i++) {
        c[i] = a[i] + b[i];
        printf("tid= %d i= %d c[i]= %f\n", tid,
i, c[i]);
    }
} /* end of parallel for */
```

- Compile-time or runtime bug?

Example: Find the bug

```
/* Declare and initialize a, b, c, chunk, ... */

#pragma omp parallel for shared(a,b,c,chunk)
private(i,tid) schedule(static,chunk) {
    tid = omp_get_thread_num();
    for (i=0; i < N; i++) {
        c[i] = a[i] + b[i];
        printf("tid= %d i= %d c[i]= %f\n", tid,
i, c[i]);
    }
} /* end of parallel for */
```

- Compilation error
- First line in parallel for must be the for loop

Example: Find the bug

```
#pragma omp parallel {  
    .....  
    #pragma omp for nowait  
        for(i=0; i<N; i++) {  
            a[i] = b[i] + c[i];  
        }  
    ..... // No barriers  
    #pragma omp for schedule (dynamic)  
        for(i=0; i<2*N; i++) {  
            c[i] = e[i] + f[i];  
        }  
    .....  
}
```

- Compile-time or runtime?

Example: Find the bug

```
#pragma omp parallel {  
.....  
#pragma omp for nowait  
    for(i=0; i<N; i++) {  
        a[i] = b[i] + c[i];  
    }  
..... // No barriers  
#pragma omp for schedule (dynamic)  
    for(i=0; i<N; i++) {  
        c[i] = e[i] + f[i];  
    }  
.....  
}
```

- Runtime error (will not segfault)
- Answer may be wrong (race)